


# Assignment 2

## Search and Optimization


CS24M109 - Pashaula Eswar Sai

CS24M123 - Gooty Bharadwaj


# Branch and Bound



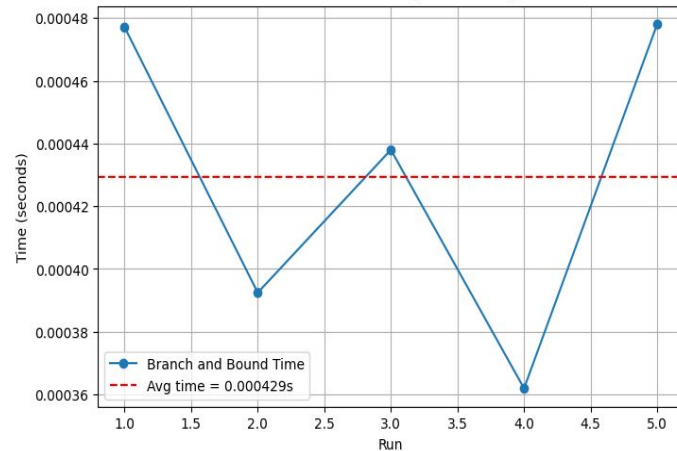
BnB: Convergence Steps (Frozen Lake)



BnB: Reward per Run (Frozen Lake)




BnB: Time to Goal (Frozen Lake)



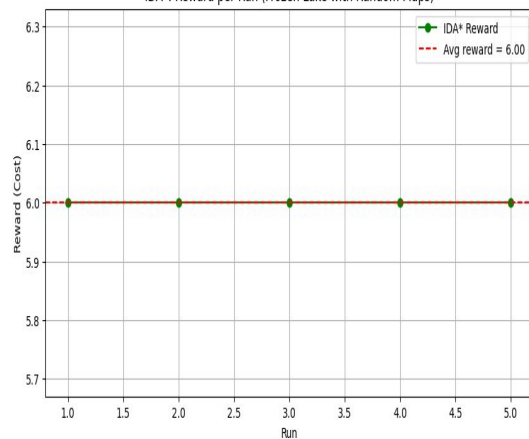
# IDA\*



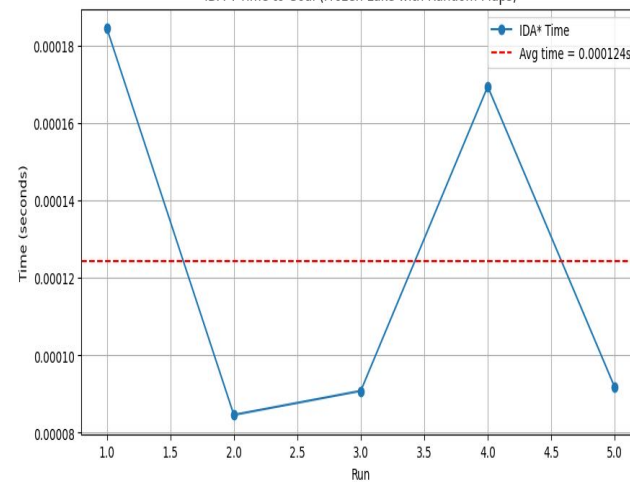
IDA\*: Convergence Steps (Frozen Lake with Random Maps)



IDA\*: Reward per Run (Frozen Lake with Random Maps)



IDA\*: Time to Goal (Frozen Lake with Random Maps)



# Heuristic Function

---

For BnB and IDA\* the Manhattan distance:

```
# Heuristic function: Manhattan Distance
def heuristic(state, goal, size=4):
    x1, y1 = state // size, state % size
    x2, y2 = goal // size, goal % size
    return abs(x1 - x2) + abs(y1 - y2)
```

For Hill Climbing and Simulated Annealing




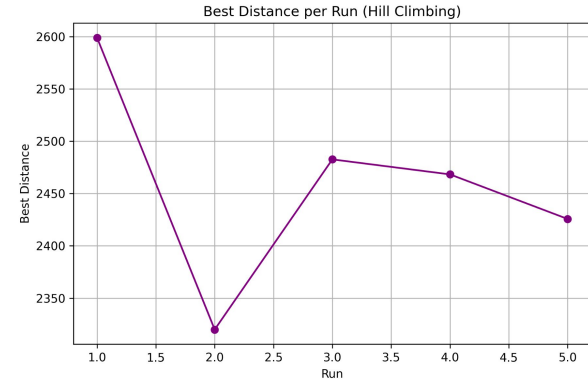
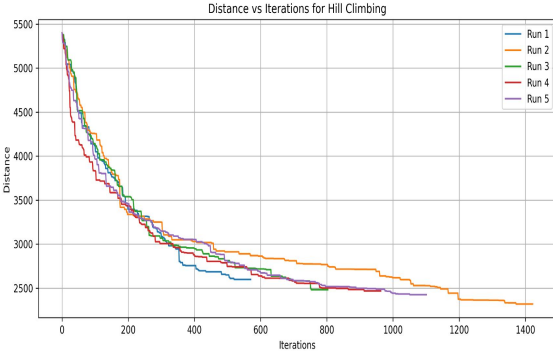
```
def total_distance(tour, cities):
    return sum(np.linalg.norm(cities[tour[i]] - cities[tour[i - 1]])) for i in range(len(tour))
```

# Observations for BnB and IDA\*

---


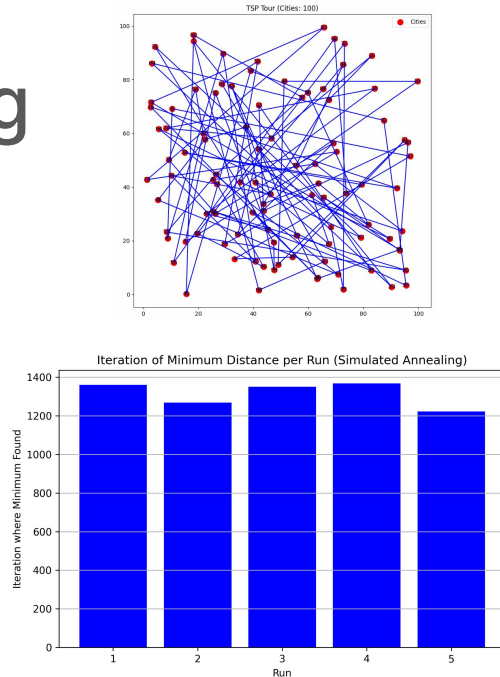

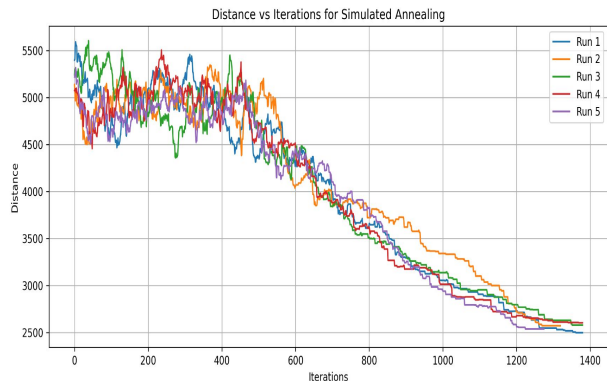
- Both IDA\* and BnB were run on Frozen Lake for a fixed start and goal point while changing the positions of obstacles in the path.
- IDA\* performed ~3.5x faster than BnB on average. This happens because IDA\* leverages heuristic-guided along with depth limits pruning large parts of the search space early.
- It avoids exploring full subtrees that exceed the thresholds.
- BnB, while using cost-based pruning, still explores more suboptimal branches before updating the best path.
- Both algorithms are complete and optimal for the frozen lake problem.
- IDA\* is better suited for Frozen Lake compared to BnB. BnB has overhead in terms of recursion and storage because of keeping track of visited states and costs
- Both algorithms found optimal solutions in all runs (Cost=6)

# Hill Climbing for Travelling Salesman



1. **Consistent downward trend** – Hill climbing always chooses better solutions, leading to steady improvement.
2. **No sudden jumps** – Only better neighbors are accepted, so there are no sharp distance drops.
3. **Early improvements** – Most of the improvement happens in the initial iterations when better neighbors are abundant.
4. **Plateaus/stagnation** – Some runs get stuck in local minima, leading to flat lines (no improvement).
5. **Different end values** – The final distance varies due to different starting points and local optima.

# Simulated Annealing



1. **High fluctuations early** – Simulated annealing accepts worse solutions at high temperatures to escape local minima.
2. **Smoother drop later** – As the temperature decreases, it becomes more selective, accepting only better or similar solutions.
3. **Sudden drops** – Occasionally, a significantly better solution is found, leading to sharp improvements.
4. **Different early paths** – Random initial solutions cause each run to explore a different trajectory at the beginning.
5. **Similar final values** – Despite randomness, the algorithm converges towards a near-optimal solution in all runs.